

ACQUISITION READINESS & VALUATION ANALYSIS

AcquisitionOS

An evidence-based internal assessment of the entire product: verified business metrics extracted from the production database, technology depth, market benchmarks, and honest valuation ranges under conservative, base, and strategic acquisition scenarios.

Internal Analysis - Prepared for the Founder

Based on direct codebase inspection and production data

Scenario ranges labeled Verified / Benchmark / Estimate / Assumption

SEPTEMBER 2026

Table of Contents

1. Executive Summary	2
2. What AcquisitionOS Actually Is	2
3. Current Product Maturity	3
4. Verified Business Metrics	4
5. Technology Assessment	4
6. Strategic Value	5
7. Current Market Benchmarks	6
8. Valuation Scenarios	6
9. Conservative Range	7
10. Base Range	7
11. Strategic Range	7
12. Key Assumptions	8
13. Value-Increasing Factors	8
14. Value-Reducing Factors	9
15. Buyer Due-Diligence Risks	9
16. Acquisition Package Checklist	10
17. What Must Be Fixed Before Selling	10
18. What Additional Data Is Required	11
19. Final Current Valuation Range	11
20. Potential Valuation After Specific Improvements	12

1. Executive Summary

AcquisitionOS is an AI-powered lead-generation and outreach platform consisting of a large Next.js codebase (1,051 TypeScript files, roughly 336,000 lines of source, 496 API routes, 104 database models) with a working authentication stack, credit system, discovery engine, AI analysis services, and an extensive dashboard. This analysis was produced by direct inspection of the current codebase and production database, and every material conclusion is labeled as verified, a market benchmark, a calculated estimate, an assumption, or unknown.

The commercial picture is difficult and this report does not soften it: the platform has **29 registered accounts, of which at most 3-4 are real persons**, and **zero verified revenue** - there are no paid payment orders in the database, no deals, no outreach messages sent, and no credit-ledger activity. The three "active" paid subscriptions in the database are QA/test fixtures, not customer purchases. On current evidence the defensible value of the business lies almost entirely in its technology, not its traction.

\$0 Verified revenue (VERIFIED)	29 Registered accounts (VERIFIED)	3-4 Real-person accounts (VERIFIED)	88 Leads in production DB (VERIFIED)
---	---	---	--

The complete valuation picture: a conservative fast sale plausibly clears **USD 10,000-25,000**; a properly marketed base-case sale of the technology asset supports roughly **USD 25,000-75,000**; and a strategic buyer who could fold the workflow/AI-lead-intelligence stack into an existing distribution channel might pay **USD 100,000-250,000**, but only under conditions that do not exist today (working pilots, documented integration value, clean full-codebase security audit). These ranges are calculated estimates anchored to micro-SaaS asset benchmarks, not guarantees. In Indian Rupees (assumption: 1 USD = INR 84) the ranges are roughly INR 8.4-21 lakh, INR 21-63 lakh, and INR 84 lakh-2.1 crore respectively.

2. What AcquisitionOS Actually Is

AcquisitionOS is a single-tenant-to-multitenant-capable B2B SaaS aimed at agencies and service sellers: it discovers business leads from multiple sources (Google Custom Search, z-ai web search, and directory-style sources), enriches them, scores them, runs AI analysis (website quality, gap detection, outreach angles), and orchestrates outreach sequences and workflows with a credit-based monetization meter. The product ships a dashboard with 181 dashboard components, a workflow builder with a template library, messaging/outreach surfaces, an API-key system with plan entitlements, and billing integrations for Stripe and Razorpay.

Verified subsystems, from direct code and runtime inspection: OTP/magic-link/Google OAuth authentication with real Gmail SMTP delivery (all three proven end-to-end during this analysis session), session and MFA scaffolding, rate limiting, audit logging, a lead-resolution authorization layer, a website discovery-and-verification service, credit accounting, plan entitlement middleware, observability hooks, and 145 internal documentation files. The database is SQLite via Prisma - acceptable for a sandbox, but a genuine

production-readiness gap (see Section 5).

What it is not, honestly: it is not a proven revenue machine, and parts of the surface area are newer than they look. A git lineage reset in a prior session destroyed several feature commits (including an earlier pipeline UI and isolation fixes), which is why portions of this session were spent re-applying lost security work. Buyers performing due diligence on git history will see discontinuities; that must be explained rather than hidden.

3. Current Product Maturity

Classification: **late MVP / early beta** - functional, wide, and demonstrably improvable, but not yet commercially deployable. The evidence for "functional" is strong: the auth stack delivers real email; discovery jobs complete; the enrichment pipeline discovered and verified real company websites (cytecare.com, nakheel.com) during this session with evidence-based verification; the UI renders a full dashboard. The evidence against "commercially deployable" is equally strong: zero production payments, a database engine not suited to multi-writer production, a partially completed application-wide security audit, and only 12 test files across a 336k-line codebase.

Layer	State	Evidence / gap
Authentication & sessions	Production-usable	OTP, magic link, Google OAuth verified live; MFA scaffold present
Lead discovery & enrichment	Beta	Multi-source fan-out works; website verification added this session
AI analysis & research	Beta	Evidence guardrails added this session; prompt quality varies by provider
Workflow automation	Beta	Engine + templates present; UI for one flagship flow lost in lineage reset
Billing (Stripe/Razorpay)	Incomplete	Models and code exist; no live keys configured; zero payment orders
Data layer	MVP	SQLite in production path; no migration discipline before this session
Tests & CI	Weak	12 test files; no CI pipeline in repository
Observability & ops	Scaffold	Health/metrics routes exist; no external monitoring or alerting

Table 1 - Maturity assessment by layer (VERIFIED unless noted).

Known defects as of this session: an unauthenticated website-analysis endpoint (fixed), cross-account lead visibility via a fail-open org check (fixed), missing user scoping on lead creation and search (fixed), and a full 40-plus-resource isolation audit that remains pending. Technical debt is concentrated in duplicated scaffolding components, inconsistent authorization middleware adoption across the 496 routes, and sparse tests.

4. Verified Business Metrics

All figures below were queried directly from the production database (db/custom.db) on 21 September 2026. Where a metric cannot be established from data, it is explicitly marked Unknown; nothing here is invented.

Metric	Value	Status	Note
Registered users	29	VERIFIED	Includes QA/pentest fixtures
Real-person accounts	3-4	VERIFIED	Owner accounts + one 2026-06 signup
Active users (30d / 7d)	3 / 2	VERIFIED	Effectively founder + testers
Paying customers	0	VERIFIED	No paid PaymentOrder rows exist
MRR / ARR	0 / 0	VERIFIED	Revenue-multiple valuation base is nil
Active paid subscriptions	3 pro + 1 elite	VERIFIED (fixture)	Created by QA scripts, not checkout
Leads in DB	88 (29 with website)	VERIFIED	87 in "discovered" stage, 1 analyzed
Outreach messages / emails sent	0 / 0	VERIFIED	Core loop never exercised at scale
Deals / pipeline value	0 / 0	VERIFIED	No sales execution history
Credit ledger rows	0	VERIFIED	Credit system never used commercially
API keys	2	VERIFIED	Developer-access surface present
Workflow definitions / executions	1 / 2	VERIFIED	Includes this session's tests
Churn, retention, CAC, LTV	-	UNKNOWN	No cohort history to compute from
AI/API cost and infra cost	-	UNKNOWN	No cost records exported by the platform

Table 2 - Verified metrics from the production database (21 Sep 2026).

Interpretation: these numbers describe a product with real usage by its builder and effectively no external adoption. That does not make the asset worthless - it makes it an asset sale rather than a business sale, and it means any buyer thesis must rest on the technology and the buyer's own distribution, not on acquired customers or revenue. Customer concentration is total: one hundred percent of meaningful activity traces to the founder's accounts.

5. Technology Assessment

Engineering surface: 1,051 TypeScript/TSX files and roughly 336,000 lines of source, 496 API route files, 231 React components (181 dashboard plus 50 UI primitives), 104 Prisma models, 172 library modules including a dedicated AI module layer, and 145 markdown documentation files. A large share of the

component and route count is scaffolded breadth rather than depth - a buyer should value the working cores (auth, discovery, enrichment, analysis, credits, workflows, audit), not the raw line count.

Subsystem	Depth assessment	Rebuild effort (ESTIMATE)
Auth & account security	Deep: three login paths, sessions, MFA scaffold, rate limits, audit trail	4-6 engineer-weeks
Lead discovery engine	Solid: source fan-out, dedup, scraping metrics, website discovery + verification	6-8 engineer-weeks
AI analysis & research	Moderate-deep: provider failover, prompt manager, evidence guardrails, cost metering	4-6 engineer-weeks
Workflow engine	Moderate: definitions, executions, triggers, templates, DLQ, credits integration	5-7 engineer-weeks
Outreach & messaging	Broad but thin: many channels scaffolded, none exercised in production	4-6 engineer-weeks
Billing & entitlements	Partial: models + middleware present, no live gateway keys, zero transactions	3-5 engineer-weeks
Dashboard UI	Very broad, uneven: 181 components, some duplicated or orphaned scaffolding	8-12 engineer-weeks
API surface	Very broad: 496 route files with inconsistent middleware adoption	audit + 6-10 engineer-weeks

Table 3 - Major subsystems and engineering replacement perspective (ESTIMATE where noted).

At Indian market rates a like-for-like rebuild of the verified cores plausibly costs USD 60,000-180,000 in engineer time (ESTIMATE); at US senior rates the same work is USD 250,000-600,000 (ESTIMATE). Replacement cost is a supporting perspective only - a no-revenue codebase does not sell at replacement cost, and sophisticated buyers know it. Security posture after this session's fixes is meaningfully better than before (strict lead resolution, evidence guardrails), but the codebase still requires the full isolation audit before any buyer pentest. Scalability is the clearest technical ceiling: SQLite plus in-process background jobs will not survive a multi-tenant production launch without a PostgreSQL migration and a queueing layer (VERIFIED gap).

6. Strategic Value

The proprietary-feeling assets are: the website discovery-and-verification engine (candidate extraction, multi-signal company-match verification with confidence levels and evidence capture - genuinely uncommon in micro-SaaS), the credit-metered AI workflow architecture, the multi-source discovery fan-out, and the evidence-guardrail pattern that deterministically strips unsupported AI conclusions. These are reusable components a strategic buyer could embed in an existing product rather than rebuild.

Plausible strategic buyer profiles (capability match, not a claim that any specific buyer will purchase): sales-automation and lead-generation vendors wanting an AI research layer; marketing agencies wanting a white-label prospecting engine; CRM platforms wanting enrichment workflows; Indian SaaS portfolio buyers wanting a ready AI team-lead skeleton. For such buyers the value is time-to-capability: skipping three to six months of build. No legal/IP verification has been performed; ownership, licenses of third-party dependencies, and prompt/data provenance all require counsel review before any sale (flagged, not assessed).

7. Current Market Benchmarks

Benchmarks below were retrieved from current public sources in September 2026 and are market data, not AcquisitionOS-specific evidence.

Source / dataset	Benchmark
FE International, 2026, sub-USD 1M ARR SaaS	2x-3x ARR (top quartile higher)
Acquire.com biannual multiples report	~5.5x ARR average by end-2025 (down from ~7x)
Flippa micro-SaaS data, 2026	Sub-USD 1M ARR sells ~2.85x annual profit; top quartile ~6.13x
Private SaaS ranges, 2026 commentary	4x-9x ARR for healthy growth; public index fell to ~3.2x mid-2026
Software Equity Group deal volume	2,698 SaaS M&A; transactions in 2025, +28% YoY

Table 4 - 2026 SaaS acquisition benchmarks (MARKET BENCHMARK).

Every one of these multiples multiplies revenue or profit. AcquisitionOS has none, so the benchmark table bounds the revenue path (Scenario C yields zero today) and pushes the valuation exercise to asset-value reasoning: what does a working, broad, security-hardened-but-not-audited AI SaaS codebase with near-zero traction actually trade for? Marketplace observation says low four figures to low five figures for typical no-revenue micro-SaaS, with outliers above that when the buyer has an immediate strategic use (ESTIMATE anchored on the benchmarks above).

8. Valuation Scenarios

Scenario A - Technology/asset sale. The buyer acquires source code, architecture, prompts, workflows, and documentation, with no meaningful recurring revenue. Anchored to no-traction micro-SaaS marketplace outcomes and discounted replacement cost, a realistic range is USD 10,000-60,000, with disciplined negotiation and a clean audit package pushing toward the upper end (CALCULATED ESTIMATE).

Scenario B - Product acquisition (working SaaS with users). This scenario collapses toward Scenario A because the verified user base is three to four real accounts and the paying base is zero. A buyer who insists on valuing "users" here would be valuing the founder's own test accounts (VERIFIED data; therefore Scenario B

range equals Scenario A).

Scenario C - Revenue-generating SaaS. ARR is verifiably zero, so any revenue multiple applied today yields zero; this scenario becomes relevant only after real revenue exists. Illustratively (and only illustratively): at the FE International 2x-3x sub-USD 1M band, every USD 10,000 of verified ARR would add roughly USD 20,000-30,000 of enterprise value (BENCHMARK arithmetic, not a forecast).

Scenario D - Strategic acquisition. A premium is possible only where the buyer can immediately attach the product to an existing customer base and recover their integration cost within months. Conditions required: the buyer already sells to the same ICP; the codebase passes their security review (pending full audit); and at least one or two working external pilots de-risk the "does anyone want this" question. None of those conditions is met today, so the strategic range is conditional, not current (ASSUMPTION-based scenario).

9. Conservative Range

A fast, low-risk sale - priced for certainty, sold quickly, likely as-is with limited support - plausibly clears USD 10,000-25,000 (INR 8.4-21 lakh at the stated FX assumption). This is the realistic floor zone for a serious, arms-length buyer of the complete asset (code, docs, DB, and handover), and it assumes the security-audit and documentation gaps disclosed honestly during diligence (CALCULATED ESTIMATE).

10. Base Range

A reasonably marketed sale - proper deal packaging, a data room, two to three months of exposure to micro-SaaS and agency-tool buyers, and a founder available for a transition period - supports roughly USD 25,000-75,000 (INR 21-63 lakh). The upper end requires the pre-sale fixes in Section 17 to be completed and evidenced, plus at least one credible external pilot or a demonstrable waitlist to counter the zero-traction objection (CALCULATED ESTIMATE).

11. Strategic Range

A strategically aligned buyer valuing capability and time-to-market over current financials could justify USD 100,000-250,000 (INR 84 lakh-2.1 crore). This is explicitly conditional: it requires working pilots with real users, a clean application-wide security audit, PostgreSQL-grade deployment, and demonstrated integration into the buyer's distribution. Without those conditions the strategic range is theoretical; with them, the workflow/AI-research stack is the kind of component a larger platform might otherwise spend six figures and two quarters building (ASSUMPTION-based scenario).




Figure 1 - Indicative valuation ranges under the three scenarios (CALCULATED ESTIMATE).

12. Key Assumptions

FX assumption: 1 USD = INR 84 for all conversions; actual transaction FX will differ (ASSUMPTION).

Scope assumption: the sale is the complete product - source, docs, current database, and handover support - not a partial asset (ASSUMPTION).

Data assumption: the production database snapshot analyzed here is complete and unmodified; a buyer will re-run these queries (VERIFIED for this session).

Security assumption: the cross-account isolation fixes verified today hold under a full-codebase audit; until that audit runs, residual unknowns remain (ASSUMPTION over VERIFIED fixes).

Market anchoring: micro-SaaS asset benchmarks from FE International, Acquire.com, and Flippa 2025-2026 data approximate the buyer pool for this asset (MARKET BENCHMARK).

No undisclosed liabilities: no external contracts, debts, or license conflicts were found in the workspace, but external accounts (Google Cloud, Gmail, search API) sit outside the codebase and were not audited (UNKNOWN).

13. Value-Increasing Factors

Verified working AI cores: discovery, enrichment with website verification, AI analysis with evidence guardrails, and a credit meter - all demonstrated end-to-end this session (VERIFIED).

Breadth of the platform: 104 database models and 496 API routes give a buyer many entry points to productize without starting from scratch (VERIFIED breadth; depth uneven).

Real authentication stack: OTP, magic link, and Google OAuth with actual email delivery proven live today, plus sessions, MFA scaffolding, rate limiting, and audit logging (VERIFIED).

Documentation volume: 145 internal docs including handover and secrets-management notes reduce diligence friction (VERIFIED; quality varies).

Security trajectory: the isolation fixes re-applied and verified this session show the codebase can be hardened quickly when prioritized (VERIFIED fixes; audit pending).

Clean licensing baseline: standard open-source stack (Next.js, Prisma, Tailwind) with no obvious copyleft contamination in core code (VERIFIED at dependency-list level; formal scan pending).

14. Value-Reducing Factors

Zero revenue and near-zero users: the single largest discount; every revenue-multiple framework prices this at zero (VERIFIED).

Founder dependency: deployment knowledge lives in sandbox scripts and personal accounts (Gmail SMTP, Google Cloud, search API keys); no operational runbook exists for a third party (VERIFIED).

SQLite in the production path: flags scalability immaturity and implies a mandatory migration before real launch (VERIFIED).

Thin tests: 12 test files against 336k lines; a buyer must budget regression risk (VERIFIED).

Git history discontinuity: a lineage reset destroyed prior commits, including once-shipped fixes; history will not read as a clean single narrative (VERIFIED).

Uneven authorization coverage: this session fixed several fail-open patterns; the remaining 400-plus routes are audited only by category, not exhaustively (VERIFIED gap).

Scaffolding inflation: a material share of components/routes is unused or duplicated, which inflates apparent size and will read as noise to technical buyers (VERIFIED).

Payment integrations incomplete: Stripe/Razorpay code exists but has never transacted; gateway accounts are not transferable assets (VERIFIED).

15. Buyer Due-Diligence Risks

A buyer performing diligence today would find the codebase organized and documented but would stumble on operational ownership and history. The checklist below states what a buyer will ask for and the current state; "missing" items are the negotiation-leverage losses, not necessarily deal-breakers.

Item	Current state	Buyer concern level
Source code + git history	Present, but lineage discontinuity	Medium
Architecture documentation	145 docs, uneven quality	Low
Deployment runbook	Sandbox-specific scripts only	High
Secrets management	env vault script; secrets in personal accounts	High
Database migrations	Started this session (2 historical + 1 new)	Medium
Automated tests / CI	12 test files; no CI	High
Monitoring / logging / backups	In-app scaffolding; no external ops	High
Security audit	Category-level; full audit pending	High
Payment system	Code present, never transacted	Medium
Third-party API ownership	Gmail/Google Cloud/search keys personal	Medium
Legal/IP documents	None in workspace	High
Customer data & privacy posture	Almost none held; GDPR pages scaffolded	Low

Table 5 - Due-diligence readiness checklist.

16. Acquisition Package Checklist

The complete handover package a buyer should receive, with current availability: product walkthrough and feature map (docs exist; needs curation); full source repository with corrected history notes (ready); environment template with every variable documented minus secret values (partially done via the vault script); cloud and third-party account transfers - Google Cloud OAuth client, Gmail app password or a proper transactional email account, Google Custom Search key, z-ai SDK account (owner action required); database with a documented schema and the migration discipline started this session (ready); payment gateway accounts or fresh registrations (Stripe/Razorpay - none exist yet); domain and brand assets (owner inventory required); analytics exports (nothing meaningful to hand over); legal/IP assignment documents (to be drafted by counsel); and a two-to-four week founder transition window for deployment knowledge transfer (recommended, and priced into the base range).

17. What Must Be Fixed Before Selling

Priority	Action	Why it moves value
P0	Complete the application-wide authorization audit (all resource types, all 496 routes)	One more fail-open leak discovered by a buyer re-prices everything
P0	Migrate SQLite to PostgreSQL; add a real deployment runbook	Removes the "prototype data layer" objection
P1	Stand up CI with a smoke suite over the verified cores	Converts "336k lines, 12 tests" fear into evidence
P1	Move email/search/oauth to owned accounts; write the secrets runbook	Removes founder-dependency from the critical path
P1	Get 2-3 external pilots (even free) with usage data	The cheapest possible answer to zero traction
P2	Prune dead scaffolding; document the true core surface	Diligence cleanliness; shrinks the perceived-audit surface
P2	Wire one payment gateway end-to-end with a real transaction	Turns billing from "code exists" to "pipeline works"

Table 6 - Pre-sale priorities ordered by valuation impact (ESTIMATE for impact).

18. What Additional Data Is Required

Cost records: LLM/API spend per analysis and infra cost per month - required to argue unit economics (currently UNKNOWN).

Cohort data: any signup/activation/retention history once real users exist; today there is none to analyze (UNKNOWN).

Conversion evidence: discovery-to-analysis-to-outreach funnel counts from real campaigns rather than QA fixtures (UNKNOWN).

Competitive pricing evidence: what comparable AI prospecting tools charge, to justify the credit price list (UNKNOWN).

Legal inventory: IP assignments, domain ownership, third-party account ownership, and any contractual constraints (UNKNOWN).

External security report: an independent audit or pentest report would materially de-risk the buyer's own review (not yet performed).

19. Final Current Valuation Range

Direct answer to "what could AcquisitionOS realistically sell for today": **Conservative USD 10,000-25,000 (INR 8.4-21 lakh); Base USD 25,000-75,000 (INR 21-63 lakh); Strategic USD 100,000-250,000 (INR 84 lakh-2.1 crore, conditional)**. The conservative and base ranges price the technology and the option it

represents; the strategic range prices what the platform could become inside an existing distribution channel, and it requires conditions - pilots, audit, infrastructure - that do not currently hold.

These are enterprise-value framings for a 100 percent asset sale; asking price, equity value, cash at closing, earnouts, and seller financing are deal-structure variables a broker or counsel would layer on top. In the current state, an earnout tied to post-acquisition activation is the most likely structure a buyer would propose, precisely because the verified traction base is so thin (CALCULATED ESTIMATE grounded in the benchmarks of Section 7).

20. Potential Valuation After Specific Improvements

The table below shows how specific, verifiable changes would move the value range. Valuation-impact figures are directional estimates anchored to the Section 7 benchmarks - not promises - and only become "real" when the underlying metric is real.

Metric / improvement	Current state	Target	Potential valuation impact
Verified MRR	USD 0 (VERIFIED)	USD 1-2k MRR from 10-20 paying users	Moves valuation from asset-base (USD 25-75k) to revenue-base (2x-3x ARR = USD 24-72k, overlapping but bankable)
Real activated users	3-4 accounts (VERIFIED)	25-50 with weekly activity	Unlocks Scenario B credibility; +30-60% on base range
Security audit	Pending (VERIFIED gap)	Independent pass + fixes	Removes the largest diligence haircut; +10-25%
Data layer	SQLite (VERIFIED)	PostgreSQL + runbook + backups	Removes prototype objection; +10-20%
Pilot case studies	0 (VERIFIED)	2-3 documented outcomes	Supports strategic-range conversations
Founder dependency	High (VERIFIED)	Runbook + 2-4 week transition	Reduces execution risk discount

Table 7 - Improvement-to-valuation map (ESTIMATE).

Bottom line: as it stands today the product is a technically rich, commercially unproven asset whose honest base value sits in the tens of thousands of dollars, not hundreds. The fastest genuine value creators are real revenue and an independent security pass - both achievable, neither fudgeable, and both exactly what a serious buyer will check first. This report deliberately offers no sale recommendation; it documents what the evidence supports.