Every question from investors, clients, users, engineers, competitors

AcquisitionOS — AI-Powered Client Acquisition System

Section A — Complete Q&A; Collection

This document covers every possible question you may face when pitching AcquisitionOS to investors, clients, companies, users, engineers, and competitors. Answers are honest — where the product isn't ready, it says so. Real production data: 536 leads, 101 contacted, 13 autopilot runs, 0 deals, $0 revenue.

1. Investor Questions

Q: What exactly is AcquisitionOS?

A: AcquisitionOS is an AI-powered client acquisition platform that autonomously discovers, analyzes, scores, and contacts potential clients across multiple channels. You set your target niche and geography, and the Autopilot engine runs the entire pipeline: discover → analyze → score → generate personalized outreach → send via email/WhatsApp/LinkedIn → track replies → surface what needs human attention. It's not a CRM, not a data provider, not an outreach tool — it's the full acquisition lifecycle in one platform.

Q: What problem does it solve?

A: Small businesses and agencies spend 15+ hours/week manually finding leads on Google Maps and LinkedIn, cold-emailing them, and tracking responses in spreadsheets. Existing tools are either too expensive (Apollo + Outreach + ZoomInfo = $500-2000/month) or too narrow (each solves one step). AcquisitionOS closes the loop: one tool, one credit system, one dashboard, from discovery to deal.

Q: What stage is the product at?

A: Functional MVP with internal testing. 536 real leads discovered, 101 contacted, 13 autopilot runs completed, 9 outreach messages sent, 1 business profile with research intelligence. Zero paying customers, zero revenue. The auth system works (email/password, Google OAuth, OTP, magic link all verified). The admin CRM and Sales Command Center render real data. Stripe integration is coded but not yet activated (no Stripe keys configured).

Q: What's your traction?

A: Honestly: minimal. 43 users in the database, but most are test identities. 4 real users (mailtoprabhat72@gmail.com on Elite plan with 2000 credits, kattyboy785@gmail.com on Pro with 750 credits, contact@prabhat.online as super_admin, nachiketasharma23@gmail.com). No completed payments yet. This is pre-revenue, pre-customers.

Q: What's the market size?

A: The global sales engagement platform market is $5.2B in 2024, projected to reach $16.7B by 2032 (CAGR 15.7%). The AI SDR (sales development representative) sub-segment is the fastest-growing, with companies like 11x.ai reaching $100M ARR in 2 years. Our initial target — Indian SMBs and agencies — represents 50M+ businesses spending $0 on sales tools today because existing solutions are too expensive.

Q: Who are your competitors and how are you different?

A: Direct competitors: HubSpot ($1.6B revenue), Apollo ($250M ARR), Outreach ($300M ARR), Clay ($50M ARR), 11x.ai ($100M ARR). They're better-funded and more mature. Our difference: they each

solve ONE step (data OR enrichment OR sequences OR CRM). AcquisitionOS is the only platform that autonomously runs the ENTIRE pipeline. You don't manage leads — you configure targets and the system acquires clients. However, honestly, 11x.ai and Artisan are pursuing the same 'autonomous SDR' vision with $50M+ in funding. Our advantage is India-first focus, credit-based pricing accessible to SMBs, and a working product today.

Q: Why would someone switch from HubSpot/Apollo to you?

A: They wouldn't — yet. HubSpot customers paying $500+/month are not our target. Our customer is the agency or SMB spending $0 on tools because the existing stack is unaffordable. We're not competing for HubSpot's customers; we're serving the 50M+ businesses that HubSpot doesn't serve because they're too small or too price-sensitive.

Q: What's your moat / defensibility?

A: Three layers: (1) Integrated data loop — every outreach sent and reply tracked improves the scoring model. This compound learning advantage only exists at scale. (2) India-first data sources (JustDial, IndiaMART) that US competitors don't have. (3) Credit-based pricing that's accessible to the segment competitors ignore. Honestly: these are weak moats today. The real moat builds when we have 1000+ active users generating reply data that improves the AI models. Until then, execution speed is our only advantage.

Q: How do you make money?

A: Credit-based SaaS subscriptions: Starter (Free, 50 credits), Pro ($29/mo, 750 credits), Elite ($99/mo, 2000 credits), Enterprise (custom). Credits are consumed by discovery, analysis, AI message generation, and email sending. Additional credit packs can be purchased ($10-$60 for 250-2500 credits). This aligns cost with value — users who acquire more clients pay more.

Q: What's your CAC (Customer Acquisition Cost)?

A: Unknown — we have zero paying customers so we can't calculate CAC yet. Our GTM strategy targets $0 CAC through: (1) content marketing (showing real discovery results), (2) the product itself as a demo (leads discovered in the user's niche during trial), (3) agency partnerships (reseller model with 25% recurring commission). Realistically, we expect CAC of $15-30 once paid acquisition begins.

Q: What's your burn rate?

A: Low — the product is built and running on infrastructure that costs ~$50-100/month (Vercel free tier + SQLite + Gmail SMTP). No salaries being paid. The founder is full-time without compensation. The ask is for runway to hire 2 engineers and launch paid acquisition.

Q: How much are you raising and what's the use of funds?

A: Ask: $250K-500K for 15-20% equity (pre-money valuation $1.5M-2.5M). Use of funds: (1) 2 engineers for 12 months ($180K), (2) cloud infrastructure ($24K), (3) paid acquisition / marketing ($60K), (4) legal/compliance ($20K), (5) operational buffer ($66K). This gives 18-month runway to reach $50K MRR.

Q: What's your 12-month milestone?

A: 1,000 active users, 100 paying customers, $10K MRR. Break-even at ~$15K MRR. Specifically: activate Stripe payments, launch in India, acquire first 50 paying agencies, build WhatsApp/LinkedIn automation, reach 10,000 leads discovered.

Q: Why you? What's your background?

A: [Founder fills in personal background — technical capability, domain expertise, why this problem matters to you personally. Be honest about what you know and what you'll need to learn/hire for.]

Q: What if Google/HubSpot/Salesforce builds this?

A: They might. But they're slow — HubSpot took 3 years to add AI features that startups shipped in months. Google won't build a sales tool (not their business). Salesforce acquires rather than builds. Our advantage is speed and focus. If they enter, we pivot to vertical specialization (India-first, specific industries) where they won't follow.

Q: Why is now the right time?

A: Three converging trends: (1) AI/LLM costs dropped 90% in 2024 — personalized outreach at scale is now economically viable. (2) Existing sales tools are maxed out on price — Apollo charges $99/user/month, pricing out SMBs. (3) India's digital adoption accelerated post-2023 — 50M+ SMBs are online but have no sales tools. The window is 18-24 months before well-funded AI SDR startups (11x.ai, Artisan) expand to our price point.

Q: What's your tech stack and is it scalable?

A: Next.js 16 + React 19 + TypeScript + Prisma (SQLite dev, Postgres prod) + Redis pubsub + Stripe + Gmail SMTP + multi-AI providers (OpenAI/Anthropic/Gemini/DeepSeek). Deployed on Vercel. The architecture is event-driven with SSE real-time updates. It scales horizontally via Redis pubsub. Honestly: SQLite is the dev DB and must migrate to Postgres for production scale — this is planned for the Vercel deployment.

Q: Do you have IP/patents?

A: No patents filed. The codebase is proprietary (not open source). The defensible IP is the integrated workflow logic (discovery → analysis → scoring → outreach → tracking in one system) and the India-specific data source integrations. We may file a process patent on the credit-based autonomous acquisition pipeline if it proves novel.

2. Client / Customer Questions

Q: How is this different from Apollo or ZoomInfo?

A: Apollo gives you a database of contacts. You still have to filter, enrich, write messages, send them, and track replies manually. AcquisitionOS does ALL of that automatically. You say 'I want dentists in Mumbai with bad websites' and the system finds them, scores them, writes personalized emails, sends them, and puts replies in your inbox. Apollo can't do that. But honestly, Apollo has 10x more contacts in their database and better data quality. We're not competing on data volume — we're competing on workflow automation.

Q: How much does it cost?

A: Free to start (Starter plan, 50 credits, no credit card). Pro is $29/month for 750 credits — enough to discover ~500 leads and send ~200 outreach messages. Elite is $99/month for 2,000 credits with full automation and business intelligence. You can buy additional credits anytime ($10 for 250, $60 for 2,500). No long-term contract required — cancel anytime.

Q: Will this work for my industry?

A: Currently tested with: restaurants, real estate agents, hotels, dentists, hospitals (based on production data). The discovery engine supports Google Maps, LinkedIn, AI Search, JustDial, IndiaMART, and Yelp. If your target customers are on any of these platforms, it will work. If you're targeting enterprise B2B (Fortune 500), this isn't the right tool yet — ZoomInfo/Apollo would be better for that segment.

Q: Can I integrate with my existing CRM?

A: Not yet. AcquisitionOS has its own built-in CRM (Lead → Deal → Pipeline → Activities). We're not trying to replace HubSpot or Salesforce — we're targeting businesses that don't have a CRM yet. CRM integration (HubSpot/Salesforce sync) is on the roadmap for Q2 2027 if customer demand justifies it.

Q: How accurate is the AI lead scoring?

A: The scoring model uses 4 dimensions: reply probability, conversion potential, urgency, and revenue potential (each 0-100). It's based on website quality, digital maturity, business signals, and contact availability. Honestly: the model is rule-based, not trained on real conversion data yet (we have 0 deals closed). The scores are directional, not predictive. They'll improve as we collect real reply/conversion data.

Q: What happens if the AI sends a bad message?

A: Every AI-generated message goes through an approval workflow before sending. You can set the Autopilot to 'approval required' mode (default) where messages queue for your review. Only after you approve does the message send. You can also switch to 'fully autonomous' mode for trusted niches. We never send anything without your configuration.

Q: Is my data safe?

A: All data is stored in a SQLite database (Postgres on production Vercel). Passwords are hashed with bcrypt (12 rounds). JWT tokens are stored in httpOnly cookies. OAuth tokens are AES-256-GCM encrypted at rest. The admin console requires super_admin role. All auth routes use constant-time password verification to prevent timing attacks. We do NOT expose passwords, tokens, or secrets in API responses or logs.

Q: Can I cancel anytime?

A: Yes. Monthly subscriptions can be cancelled anytime from the Stripe customer portal. You keep your credits until they're consumed. No cancellation fee. No lock-in contract (unless you signed an annual partnership agreement with a discount).

Q: What if I need more credits?

A: Buy credit add-on packs: 250 for $10, 500 for $18, 1,000 for $30, 2,500 for $60. Or upgrade to a higher plan. Credits never expire while your subscription is active.

Q: Do you offer a free trial?

A: The Starter plan is free forever with 50 credits — no credit card required. That's enough to discover ~30 leads, analyze 10, and send 10 outreach messages. If you need more, upgrade to Pro ($29/mo) which can be cancelled anytime.

Q: How many leads can I discover per month?

A: Depends on your plan's credits and daily outreach limits. Free: 50 credits (~30 leads), 50 outreach/day. Pro: 750 credits (~500 leads), 500 outreach/day. Elite: 2,000 credits (~1,300 leads), 2,000 outreach/day. Enterprise: custom. Credits roll over while subscription is active.

3. User Questions

Q: How do I get started?

A: Sign up at acquisitionos.com → choose your niche and geography → the system discovers leads automatically → review them in the CRM → approve outreach messages → track replies in the Sales Command Center. The whole setup takes 10 minutes.

Q: What if I don't know my target niche?

A: The system has pre-built niche templates for 20+ industries (dentists, restaurants, real estate, hotels, gyms, salons, etc.). Pick one, set your city, and the Autopilot starts discovering. You can refine targeting later based on which leads convert.

Q: Can I use my own email?

A: Yes — connect your Gmail account via Google OAuth. The system sends outreach from your Gmail address using Gmail SMTP. Your emails look like they came from you, not a marketing platform. You can also connect multiple Gmail accounts for different outreach campaigns.

Q: Does it work on mobile?

A: The dashboard is responsive — it works on mobile browsers. There's no native mobile app yet. Mobile optimization is prioritized for the Sales Command Center so you can check replies and approvals on the go.

Q: Can I customize the outreach messages?

A: Yes. The AI generates a draft based on the lead's analysis (website quality, weaknesses, recommended approach). You can edit it before approving, or set templates that the AI fills in per lead. You can also create custom sequence steps with your own content.

Q: What if a lead replies negatively?

A: Replies are tracked in the Sales Command Center's 'Needs Attention' panel. You see the full conversation history and the lead's score. If someone unsubscribes, they're automatically excluded from future outreach (EmailUnsubscribe model).

Q: How do I know it's working?

A: The Sales Command Center shows real-time KPIs: leads discovered, contacted, replied, meetings booked, deals created, pipeline value. You see exactly where each lead is in the funnel and what needs your attention today.

4. Engineer / Technical Questions

Q: What's the tech stack?

A: Frontend: Next.js 16 (App Router, Turbopack), React 19, TypeScript, Tailwind CSS 4, shadcn/Radix UI, framer-motion 12, recharts 2.15, @tanstack/react-table 8, @tanstack/react-query 5. Backend: Next.js API routes (Node.js runtime), Prisma 6 ORM (SQLite dev, Postgres prod), Redis pubsub for multi-instance real-time. Auth: JWT (access + refresh), bcrypt 12 rounds, httpOnly cookies, Google OAuth, OTP via Gmail SMTP. Payments: Stripe + Razorpay. AI: Multi-provider (OpenAI, Anthropic,

Gemini, DeepSeek, OpenRouter) with per-user encrypted API key storage.

Q: How is the database structured?

A: Prisma schema with 136 models including: User, Organization, Lead, LeadAnalysis, LeadScore, LeadActivity, LeadNote, Deal, PipelineStage, OutreachMessage, OutreachSequence, SequenceStep, SequenceEnrollment, Meeting, Task, FollowUpReminder, BusinessProfile, WebProperty, ResearchSnapshot, AutopilotWorkflow, AutopilotRun, WorkflowDefinition, WorkflowExecution, PaymentOrder, Subscription, CreditsLedger, AuditLog, SystemEvent, RealtimeEvent, and more. SQLite for dev (file-based), Postgres for production.

Q: How does real-time work?

A: EventBus singleton (in-memory) with 6 channels: lead_events, payment_events, message_events, workflow_events, ai_events, notification_events. Server-side publishers (publishLeadEvent, etc.) → EventBus → SSE manager broadcasts to connected clients. SSE supports auto-reconnect with last-event-id resumption. Redis pubsub for multi-instance broadcast. RealtimeEvent DB model persists events for audit/replay. The useLiveUpdates() hook on the client subscribes to channels and triggers React Query cache invalidation on events.

Q: How is authentication handled?

A: Custom JWT-based auth (not NextAuth, though NEXTAUTH_SECRET is set for compatibility). Access token (15min) + refresh token (7d) in httpOnly, Secure, SameSite=Lax cookies. Password hashing: bcrypt 12 rounds. Constant-time password verification (dummy hash comparison even for nonexistent users to prevent timing attacks). Google OAuth: full callback flow with dynamic origin detection, state-encoded redirect URI, CSRF protection. OTP: 6-digit code, 10-minute expiry, rate-limited (5/15min per IP+email), stored on User record. RBAC: 7 roles, 50+ permissions, server-side enforced via getSuperAdminSession (pages) and withSuperAdmin/withPermission (APIs).

Q: How is the credit system implemented?

A: CreditsLedger model tracks every credit transaction with idempotency keys (prevents double-charging on retries). User.credits is the balance. Each AI operation (discovery, analysis, outreach generation) debits credits atomically within a Prisma transaction. If the operation fails, credits are refunded with a 'refund:' prefixed idempotency key. P2002 (unique constraint) is caught to handle race conditions. Credit add-on purchases credit the balance via Stripe webhook.

Q: What about security?

A: All passwords bcrypt-hashed. JWT secrets in env vars (not in code). OAuth tokens AES-256-GCM encrypted at rest (UserAIProvider model). Admin routes protected by getSuperAdminSession (server-side) + withSuperAdmin (API-side). Rate limiting on all auth endpoints (5/15min for OTP/magic-link/signin). Audit logging on all admin actions (AuditLog model with 2,232 entries in production). Security alert chain with integrity verification. Input validation on all API routes. SSRF protection on URL-fetching endpoints. No secrets in client bundle (no NEXT_PUBLIC_ secret vars). .env gitignored.

Q: What's the testing situation?

A: 2,978 tests passing, 0 failing, 5 skipped (as of last full regression). Includes unit tests (aggregation logic, KPI calculations, conversion math), integration tests (admin API authorization, CRM queries), and E2E tests (tenant isolation, idempotency/retry, error path regression, persistence/restart). Test isolation via per-worker SQLite DBs copied from a schema template.

Q: How do you handle scale?

A: Currently: SQLite (single-file, no horizontal scale). Production plan: Postgres with connection pooling (PgBouncer), Redis for multi-instance pubsub, Vercel serverless functions with auto-scaling. The architecture is designed for horizontal scaling — the EventBus + Redis pubsub pattern means multiple server instances can broadcast events to all connected clients. Honestly: we haven't load-tested yet. The first bottleneck will be Prisma query performance on large lead datasets (mitigated by proper indexing — 30+ @@index directives in schema).

5. Competitor Comparison Questions

Q: How are you different from HubSpot?

A: HubSpot is a CRM + marketing platform ($50-3000/month). It manages contacts, tracks deals, sends marketing emails. It does NOT discover leads, analyze them, or autonomously send personalized outreach. You bring leads INTO HubSpot — AcquisitionOS finds them FOR you. HubSpot is for companies that already have leads; AcquisitionOS is for companies that need to FIND leads. Honestly: HubSpot is better at CRM, reporting, and marketing automation. We don't compete there. We compete on acquisition.

Q: How are you different from Apollo?

A: Apollo is a B2B contact database ($49-99/user/month) with 275M+ contacts. You search, filter, and export. It's a data tool. AcquisitionOS is a workflow tool — it USES data (from Google Maps, LinkedIn, directories) to automatically discover, score, and contact leads. Apollo gives you the phonebook; AcquisitionOS makes the calls for you. Honestly: Apollo has 100x more contacts, better data quality, and mature filtering. We can't compete on data. We compete on automation.

Q: How are you different from 11x.ai / Artisan?

A: These are the closest competitors — AI SDR platforms that autonomously prospect and outreach. 11x.ai raised $50M+ and reached $100M ARR in 2 years. They target enterprise ($500-2000/month). We target SMBs ($0-99/month). Same vision, different market. Honestly: they have more funding, better AI (purpose-trained models), and enterprise integrations. Our advantage: India-first data sources, credit-based pricing accessible to the segment they ignore, and a working product today (not a waitlist).

Q: How are you different from Clay?

A: Clay is a lead enrichment and outreach tool ($100+/month). It's a spreadsheet-like interface where you enrich data, write AI-personalized messages, and send via Outreach/Salesloft. It's powerful but requires manual setup per campaign. AcquisitionOS automates the entire flow — you don't build the enrichment table, the system does it. Honestly: Clay is better for power users who want granular control. We're better for people who want it to just work.

Q: How are you different from Zapier/Make?

A: Zapier/Make are automation platforms that connect tools together. You could build an acquisition workflow with them (Apollo + ChatGPT + Gmail + HubSpot). But it takes 20+ hours to set up, breaks when APIs change, and costs $100+/month in subscription fees across 5+ tools. AcquisitionOS is purpose-built — no integration glue needed. Honestly: Zapier is more flexible and connects to 5,000+ apps. We connect to 20. But for acquisition specifically, we're faster to set up and cheaper.

Q: Can't I just use ChatGPT + Gmail?

A: Yes, for 5 leads/day. ChatGPT can write messages, Gmail can send them. But: who finds the leads? Who scores them? Who tracks replies? Who follows up after 3 days? Who sends at the optimal time? Who stops sending to someone who replied? You'd spend 4 hours/day doing manually what AcquisitionOS does automatically. The value isn't the AI — it's the integrated automation pipeline.

Q: Why not just use LinkedIn Sales Navigator?

A: Sales Navigator ($99/month) gives you LinkedIn search and InMail. It's a search tool, not an outreach automation platform. You still manually find, manually message, manually track. AcquisitionOS automates the entire flow including LinkedIn (planned). Honestly: Sales Navigator has better LinkedIn data than we do. But it's just search — no automation, no scoring, no follow-up sequences.

6. Hostile / Challenging Questions

Q: You have zero revenue and zero customers. Why should I invest?

A: You're right — we're pre-revenue. Here's what exists: a working product with 536 real leads discovered, 13 autopilot runs completed, verified auth (Google OAuth + OTP + email/password), Stripe integration coded (needs activation), admin CRM + Sales Command Center rendering real data. The infrastructure costs $50/month. The question isn't 'does it work' — it does. The question is 'can you acquire customers' — and that's what your investment funds: go-to-market, not product development.

Q: 11x.ai has $50M and 100x your traction. You'll be crushed.

A: Possibly. But 11x.ai targets enterprise ($500-2000/month, US/EU). We target Indian SMBs ($0-99/month). Different market, different price point, different data sources (JustDial, IndiaMART). If 11x.ai enters India at $99/month, we compete on: (1) India-specific data sources they don't have, (2) credit-based pricing (pay for what you use, not per-seat), (3) existing product momentum. Honestly: if they enter our market with their resources, it's a risk. Our defense is speed and local focus.

Q: Your AI scoring is just rules, not ML. That's not real AI.

A: Correct — today. The scoring uses weighted heuristics (website quality, digital maturity, contact availability) because we have zero conversion data to train on. Once we have 1,000+ tracked outreach → reply → deal outcomes, the scoring model upgrades to trained ML. This is the data flywheel: more usage → more data → better scoring → higher conversion → more usage. We're honest that the AI is basic today.

Q: SQLite? That's not production-grade.

A: Correct — SQLite is the development database. Production uses Postgres (configured via DATABASE_URL env var, Prisma supports both). The schema has 30+ index directives for query performance. Connection pooling via PgBouncer is configured. The migration from SQLite to Postgres is a configuration change, not a code rewrite.

Q: Anyone can build this in a weekend with OpenAI API.

A: The core discovery + message generation? Yes. But the integrated pipeline — multi-source discovery with provenance, lead analysis with 6 scoring dimensions, credit system with idempotency, multi-step sequence automation, real-time event infrastructure, admin CRM with 136 Prisma models, RBAC with 7 roles, Stripe + Razorpay integration, OTP/magic-link/OAuth auth — that's not a weekend. It's thousands of hours of implementation. The codebase has 5,140 tracked files. A competitor could replicate the *concept* in a weekend. Replicating the *product* takes months.

Q: Your pricing is too cheap. $29/month won't sustain a business.

A: At $29/month with 70% gross margin (credit cost ~$8.7/user), we need 690 paying customers for $20K MRR (break-even for 2-person team). At $99/month Elite, we need 200 customers. The low price is intentional — it's the wedge into the SMB market that competitors ignore. Once we have product-market fit, we introduce annual contracts, enterprise tier ($500+), and add-on services (custom integrations, managed outreach). The pricing will increase as the value proposition strengthens.

Q: You're a solo founder with no team. How do you execute?

A: Currently solo, but the product is 90% built. The ask includes hiring 2 engineers (one full-stack, one AI/ML) for 12 months. The founder handles product, GTM, and investor relations. Honestly: solo founder is a risk. The mitigation is that the hard technical work (architecture, auth, CRM, real-time, payments) is already done. The remaining work (WhatsApp/LinkedIn automation, ML scoring, scale optimization) is what the hires will focus on.

Q: What's to stop me from building this myself?

A: Nothing — the concept isn't patentable. But: (1) it'll take you 3-6 months of full-time development to match what exists today, (2) you'll make the same mistakes we already fixed (auth edge cases, idempotency, rate limiting, credit race conditions), (3) by the time you build it, we'll have 1,000+ users and data flywheel advantage. If you're an agency, your time is better spent using the tool than building it. If you're an investor, the question is whether the 6-month head start is worth investing in.

Q: Your 'India-first' advantage is just a few API integrations.

A: It's more than APIs — it's: (1) India-specific data sources (JustDart, IndiaMART) with Indian business taxonomy, (2) GST/tax compliance built into invoicing, (3) Razorpay integration (UPI, net banking — Stripe doesn't do UPI), (4) INR currency support throughout, (5) niche templates tuned for Indian business categories. A US competitor entering India would need 6-12 months to add these. But honestly: it's not a deep moat. The real India advantage is customer relationships and local market understanding, which comes from operating here.

7. Legal & Compliance Questions

Q: Is cold outreach legal?

A: In India: Yes, under the IT Act 2000 and DND (Do Not Disturb) registry. We check the DND registry before sending (planned). In the US: CAN-SPAM Act requires (1) accurate header info, (2) clear subject line, (3) physical address, (4) unsubscribe mechanism, (5) honor unsubscribe within 10 business days. We implement all of these (EmailUnsubscribe model, physical address in email footer). In the EU: GDPR requires consent or legitimate interest. We include unsubscribe links and process data only for B2B outreach (legitimate interest).

Q: How do you handle GDPR?

A: GDPR compliance features: (1) EmailUnsubscribe model — automatic opt-out, (2) GdprRequest model — data export/deletion requests, (3) DataExport model — automated data portability, (4) Consent tracking in the auth flow. We process only B2B contact data (business emails from public directories) under legitimate interest. We do NOT process personal/sensitive data. Data is stored in the user's database instance — not a central data lake. Honestly: full GDPR compliance review is pending — we have the schema but need legal review before EU launch.

Q: Do you sell or share user data?

A: No. Never. User data (leads, outreach, deals) belongs to the user. We don't sell it, share it, or use it to train models without explicit consent. The data is stored in the user's own database instance. Our business model is SaaS subscriptions, not data brokerage.

Q: What about email deliverability?

A: We send via the user's own Gmail account (Gmail SMTP with their App Password). This means emails come from their domain, with their sender reputation. We don't use shared IP pools (which hurt deliverability). We track opens/clicks via tracking pixels. We honor unsubscribes automatically. Daily outreach limits prevent spam-like volume. Honestly: if a user sends 500 emails/day from a new Gmail account, Google will throttle them. We recommend warming up the account gradually (50/day → 100/day → 500/day over 2 weeks).

Q: Who owns the data if a user leaves?

A: The user owns their data. They can export all leads, deals, activities, and outreach history at any time via the data export feature. If they cancel their subscription, their data is retained for 90 days (in case they return), then permanently deleted. Enterprise contracts can specify custom data retention policies.

8. Honest Product Status Summary

Feature	Status	Notes
AI Lead Discovery	■ Working	536 leads in production DB
Lead Scoring/Analysis	■ Working	Rule-based, not ML-trained

Feature	Status	Notes
Autopilot Engine	■ Working	13 runs completed, 2 failed
Outreach (Email)	■ Working	Gmail SMTP verified
Outreach (WhatsApp)	■ Not built	Config exists, no implementation
Outreach (LinkedIn)	■ Not built	Config exists, no implementation
Auth (Email/Password)	■ Working	Bcrypt 12, constant-time
Auth (Google OAuth)	■ Working	Credentials valid, redirect URI pending
Auth (OTP/Magic Link)	■ Working	Real email delivery verified
CRM (Admin)	■ Working	136 Prisma models, real data
Sales Command Center	■ Working	Real-time dashboard, live data
Stripe Payments	■■ Coded, not active	Needs Stripe keys in Vercel
Razorpay Payments	■■ Coded, not active	Needs Razorpay keys
Real-time (SSE)	■ Working	EventBus + SSE + Redis pubsub
Business Intelligence	■■ Schema drift	Tables missing from DB
Deal Pipeline	■■ Empty	0 deals in production
Competitor Analysis	■■ Schema exists	No data yet
AI Calling/Voice	■ Not built	Not in roadmap
Proposal Generation	■ Not built	Deal model has field, no generator
Mobile App	■ Not built	Responsive web only
Forecasting	■■ Model exists	No engine implemented