

Table of Contents

1. Executive Summary	1
2. What AcquisitionOS Actually Is	2
3. Current Product Maturity	3
4. Verified Business Metrics	4
5. Technology Assessment	6
6. Strategic Value	8
7. Current Market Benchmarks (2026)	8
8. Valuation Scenarios	9
9. Conservative Range	11
10. Base Range	11
11. Strategic Range	11
12. Key Assumptions and Evidence Classification	12
13. Value-Increasing Factors	12
14. Value-Reducing Factors	13
15. Buyer Due-Diligence Risks (incl. Founder Dependency & DD Readiness)	14
16. Acquisition Package Checklist	15
17. What Must Be Fixed Before Selling (incl. Exit Scorecard)	17
18. What Additional Data Is Required	18
19. Final Current Valuation Range	19
20. Potential Valuation After Specific Improvements	20

1. Executive Summary

This analysis answers one question with evidence instead of hope: if the entire AcquisitionOS product were put up for sale today, what could a real buyer reasonably pay for it? The answer below is built exclusively from direct inspection of the current workspace codebase — every route file, the Prisma schema, the test suite, the runtime database, environment configuration, git history, and documentation — and is benchmarked against published 2026 SaaS transaction data. Nothing in this report is based on the feature list, the development effort invested, or what the product could become with funding and customers it does not have.

The central findings are uncomfortable but clear. AcquisitionOS is a technically substantial product: a single Next.js 16 modular monolith with 505 API route files, roughly 258 service modules, 105 database models, and about 348,000 lines of TypeScript covering the full lead-acquisition lifecycle from discovery to outreach, meetings, billing, and an API platform. It is also a pre-commercial product with no verifiable revenue, no verifiable paying customers, a development-only database containing test accounts, unconfigured payment providers, 32 failing tests out of 824, roughly 114 hollow test skeletons that pass without asserting anything, no license file, and a git history squashed to 12 commits. The current deployment runs in a development sandbox on SQLite; no production deployment, production database, or operational telemetry was found in evidence.

\$0 Verified MRR / ARR [V]	0 Verified paying customers [V]	29 DB users — overwhelmingly test accounts [V]	\$5K–\$250K Evidence-supported value range today [CE]
--	---	--	---

Under those conditions, market evidence for pre-revenue software assets (2026) suggests most sell for \$500–\$25,000 [MB]. AcquisitionOS sits above the median pre-revenue asset on depth, breadth, and documentation, which supports a base-case asset value of roughly \$25,000–\$60,000 (₹24 lakh – ₹58 lakh) in a well-run sale, a conservative floor of \$5,000–\$15,000 in a fast sale, and a strategic ceiling of \$75,000–\$250,000 only if an aligned buyer with an existing customer base and distribution channel values the technology beyond its financials, and only after specific cleanup listed in this report. Every range assumes the buyer acquires code and architecture — not customers, revenue, or team. Sections 8–11 derive these numbers; Section 17 states the final answer; Section 19 separates verified fact from estimate throughout.

2. What AcquisitionOS Actually Is

AcquisitionOS (package name "vantage" v0.2.0 — a leftover from an earlier product iteration [V]) is an AI-powered B2B client-acquisition SaaS. Its verified core loop: discover businesses via search APIs, score their websites, research them with LLMs, compute four predictive scores (reply, conversion, urgency, revenue potential), generate personalized outreach, send it over email with open/click/reply tracking, detect buying signals in replies, orchestrate meetings through Google Calendar with a human-approval autonomy mode, and manage the resulting pipeline. A credits system meters every AI action; subscriptions (Free/Pro/Elite) with Stripe and Razorpay billing are fully coded; a Developer Access API-key platform, workflow automation engine, competitor intelligence, RAG with vector search, and multi-channel notifications (in-app, email, Telegram, WhatsApp, web push) round out the surface.

Property	Verified value [V]	Evidence
Architecture	Modular monolith — one Next.js 16 process, one database	package.json; src tree; docs/02-architecture/SYSTEM-ARCHITECTURE.md
API surface	505 route.ts files in ~40 domains	rg --files src/app/api -g route.ts
Business logic	~258 service modules in src/lib	file count, verified 2026-09-21
Data model	105 Prisma models on SQLite (db/custom.db)	prisma/schema.prisma
Frontend	Client-rendered SPA behind AuthGate; 247 components; React 19, Tailwind 4	src/app/page.tsx; component counts
Auth	Custom JWT sessions (not NextAuth runtime): email/OTP, magic link, Google OAuth, TOTP MFA, revocable sessions, RBAC	src/lib/auth.ts; /api/auth/* routes
AI stack	Z-AI SDK primary + OpenAI-compatible fallback; per-call cost metering; credit enforcement; RAG + vector search	src/lib/ai/ai-provider.ts; ai-cost-tracker.ts
Background work	No queue/worker — in-process execution + 12 HTTP cron endpoints (CRON_SECRET)	src/app/api/cron/*; ADR-011
Realtime	WebSocket (/api/ws) + SSE (/api/events/*), in-process event bus; optional Redis (unconfigured)	realtime-event-bus.ts; .env
Payments	Stripe + Razorpay fully coded (webhooks, invoices, refunds, GST, credits) but unconfigured — no keys in env	src/lib/payments/*; .env inspection
Code volume	~348,000 lines of TS/TSX in src/ (incl. AI-generated density)	wc -l over src tree

Table 1 — The product as verified by code inspection (not by documentation claims).

Two verification gaps matter for a buyer's understanding of the product. First, the production deployment story is aspirational: a reduced 6-model PostgreSQL schema (prisma/schema.production.prisma) exists for a

Supabase/Vercel target, but it covers a fraction of the 105-model dev schema, and no production database or deployment was found in evidence. Second, several headline integrations are credential-gated shells: the payment, WhatsApp, Telegram, and push channels are implemented end-to-end in code but have never processed a real transaction in this environment, because their keys are absent. A buyer is therefore acquiring working code, not working businesses plumbing.

3. Current Product Maturity

Classification: **advanced pre-commercial MVP — between Beta and Production-ready on code quality, but not commercially deployed, and not acquisition-ready as a business.** The label is justified per dimension below, because a single word would hide the asymmetry: the software is far more mature than the business around it.

Stage	Verdict	Evidence for the verdict
MVP	Exceeded	The core loop (discover → score → outreach → reply → meeting) is complete and deep, with 505 routes vs a typical MVP's dozens.
Beta	Met (software)	Feature suite docs claim 88 working / 27 partial / 3 planned features; auth, credits, workflows, and admin are implemented with real test coverage in unit/integration suites.
Production-ready	Not met operationally	SQLite in current runtime; middleware disabled; 32 failing tests; no configured payment keys; no external monitoring (Sentry/OTLP unset); no production deploy found in evidence.
Commercially deployable	Not met	No verifiable customers, revenue, or billing runs; onboarding-to-payment loop unexercised with real money.
Enterprise-ready	Not met	No SSO/SCIM, no org-level audit certifications, no SOC2/ISO artifacts, single-process scaling limits documented in docs/technical/SCALABILITY-PLAN.md.
Acquisition-ready	Not met — cleanup required	No LICENSE file, squashed git history (12 commits), fictional metrics document present in repo, legacy contradictory docs (now bannerred), no financial records.

Table 2 — Maturity classification by stage, with the evidence that forces each verdict.

3.1 Unfinished and partially implemented features

- Payments: Stripe/Razorpay flows coded but never executed end-to-end with live keys — checkout fails fast by design until configured [V].
- Middleware: src/middleware.ts disabled in this deployment (middleware.ts.disabled); edge-level checks compensated per-route [V].
- CRM sync: meeting/crm-sync.ts is a scaffold; HubSpot/Salesforce sync not implemented [V].

- Schema drift: code references fields missing from Prisma models (schema-gap-report.json); production schema covers 6 of 105 models [V].
- Tests: tests/e2e/* and tests/load/* are empty skeletons — 114 test cases with zero assertions that pass trivially [V].
- A/B testing for outreach, and multi-replica horizontal scaling: documented as planned, not implemented [V].

3.2 Technical debt, security and scalability risks

The most consequential debt is the SQLite-to-PostgreSQL gap: the entire schema, indexes, and the production variant exist, but the migration has not been run or validated at full scale, and three real migrations exist against a mostly db-push-driven workflow. Security work is genuinely present — rate limiting, CSRF, input validation, webhook signature verification, MFA, session revocation, device fingerprinting, and two internal penetration-test reports (which honestly record that only 12 of 492 routes were rate-limited at test time) — but the cross-account isolation audit that this product's history demanded was only partially completed, and a buyer will treat any residual tenancy leakage as a blocking defect. Scalability is bounded by design: one process, in-memory caching, in-process SSE/WebSocket affinity, and no queue; the repo's own scalability plan documents these limits honestly.

4. Verified Business Metrics

The only database available for inspection is the development SQLite file (db/custom.db, ~4.4 MB). No production database, analytics export, billing dashboard, or accounting record exists in the workspace. Metrics below are therefore classified as Verified (read directly from the DB or code), Estimated (calculated with stated assumptions), or Unavailable (no data source exists). Do not mistake Verified dev-database numbers for business traction: they verify what the software recorded during testing, not commercial activity.

Metric	Value	Class	Source / note
Registered users	29 in dev DB — overwhelmingly synthetic test accounts	Verified (as dev data)	users table: test@example.com, etheral test addresses, plus 2–3 manual test logins
Active users	Not measurable — no telemetry/analytics pipeline in production	Unavailable	No analytics export exists
Paying users	0 verifiable	Verified (absence)	PaymentOrder history absent; 2 invoices exist from test flows
MRR / ARR / revenue	\$0 verifiable	Verified (absence)	No live Stripe/Razorpay keys configured; no billing records
Churn / retention / conversion / CAC / LTV	Unavailable	Unavailable	No cohort or funnel data recorded
Leads created	89 rows in dev DB	Verified (as dev data)	leads table
Deals closed	0 rows	Verified (absence)	deals table empty
Subscriptions	27 rows: 23 trial, 4 marked active during test sessions	Verified (as dev data)	subscriptions table
Credit consumption / AI cost	Mechanism exists (AiCostRecord, CreditsLedger); totals are dev-only	Unavailable commercially	ai_cost_records + credits_ledger in dev DB
Infrastructure cost	Sandbox-era only; no production bill exists	Unavailable	No cloud account evidence in workspace
Customer concentration	N/A — no customers	Verified (absence)	—

Table 3 — Business metrics with evidence classification. [V] rows verify software records, not commercial traction.

The honest headline: **every financial metric that drives a revenue-multiple valuation is unavailable or zero**. Any document claiming otherwise about this product — including the fictional interview-preparation document found in the repository (“₹4.2M MRR, 3,200 organizations”), which is now labeled as fiction — has no basis in the codebase or database.

5. Technology Assessment

Judged independently of revenue, the technology is the strongest asset in the sale. The system implements a genuinely long automation chain — discovery, website scoring, AI research, four-factor scoring with explanations, outreach generation, multi-channel sending, reply intelligence, hot-lead detection, sequence enrollment, meeting autonomy with human approval, workflow engine with dead-letter handling, credits metering, and a Developer Access API — as coherent service modules rather than one tangled script. Internal modularity is real: thin routes delegate to src/lib services (a "thin routes / fat services" pattern), and domain folders (ai/, lead-discovery/, meeting/, payments/, security/, notification-channels/) mark plausible extraction boundaries.

5.1 Engineering effort represented (supporting perspective, not valuation)

Naively multiplying 348K lines by salary produces absurd numbers, because a large share of the code is AI-generated, repetitive, or scaffolding. A defensible replacement estimate works from subsystems: a skilled 2–3 person team rebuilding functional equivalence of the core loop, auth, credits, workflows, and API platform would need roughly 12–24 engineer-months; full parity of all 505 routes and auxiliary systems plausibly reaches 30–50 engineer-months. At blended fully-loaded rates of \$8K–\$15K per engineer-month, that frames a replacement cost of roughly **\$150K–\$400K for the functional core, up to ~\$600K for full parity** [CE]. This is what the code would cost to rebuild — it is not what the asset sells for, because replacement cost ignores market demand, revenue, and risk. Buyers discount it heavily; it serves only as a negotiation anchor and an insurance valuation.

5.2 Subsystem depth ratings

Subsystem	Depth (0-5)	Notes
Lead discovery + research	4	Google Custom Search + SerpAPI, website scoring with proxy rotation, AI company research; source adapters registered
AI analysis + scoring	4	Four composite scores + reasoning, prompt manager, cost tracking, enforcement before calls
Outreach + sequences	4	Generation, scheduling, open/click tracking, bounce intelligence, reply classification, autonomous SDR loop
Workflow engine	3.5	Definitions, executor, triggers, DLQ, replay, templates, metrics — in-process (no durable queue)
Meetings orchestration	4	Intent detection, slot search, approvals, Google Meet links, reminders, post-meeting AI notes
Billing + credits	3.5	Two providers, idempotent webhooks, invoices/PDF, dunning, GST — unconfigured in practice
Developer Access API	3	API keys (hashed), rotation, revocation, usage tracking, scoped endpoints
Security	3.5	MFA, session revocation, device fingerprint, rate limiting, CSRF/CORS/headers, 2 pentest reports
Observability	2.5	In-process OTel hooks, health endpoints, loggers — no external backend configured
Tests	2.5	Real unit/integration/security suites (678 passing, 1,900+ assertions); e2e/load hollow; 32 failures
Documentation	4	Current 2026-09 suite: architecture, ADRs, API reference, runbooks, KT docs; legacy contradictions now bannerred

Table 4 — Subsystem depth from code inspection. Ratings are [CE] (assessed estimates), grounded in cited files.

Deployment readiness is the weakest area: the verified runtime is a sandbox process on SQLite with a keepalive script. Multi-cloud readiness is documentation-level — deploy/k8s, terraform, render artifacts describe the abandoned FastAPI/Celery generation (now explicitly marked historical), while deploy/railway (single Docker web service) matches the current architecture and is the recommended target in the knowledge-transfer docs.

6. Strategic Value

Strategic value exists only where a buyer can plug an asset into something they already have. The capabilities below are verified in code; whether any specific buyer class would pay for them is a hypothesis to be tested in a real process — nothing here claims a named buyer will acquire the product.

Buyer type	Capability they could obtain	Why it might matter to them
Lead-gen / sales-intelligence companies	Discovery → research → scoring → outreach chain with reply intelligence	Could accelerate their roadmap by months versus building in-house; scoring logic is tunable IP
Sales-automation / AI-SDR startups	Autonomous SDR pipeline, sequence engine, hot-lead detection, meeting autonomy with approvals	Reference implementation of the full autonomous loop including human-in-the-loop controls
CRM / marketing platforms	Credits metering, entitlements, multi-provider billing scaffolding, Developer Access API	Proven metering + API-key monetization plumbing they may lack
Agencies (scaled)	White-label/org branding, multi-channel notifications, pipeline analytics	Productized client-acquisition system they could operate rather than build
PE / search-fund operators	A complete, documented codebase with KT docs and honest feature-status matrix	Lower technical DD friction; a base to roll other lead-gen acquisitions into
SaaS portfolio buyers	Workflow engine + API platform + credits as reusable infrastructure	Cross-portfolio reuse of metering and automation layers

Table 5 — Potential strategic capabilities by buyer type (hypotheses, not claims of intent).

The strongest single strategic object is the end-to-end autonomy architecture: credits-enforced AI actions, an event bus feeding SSE/WebSocket surfaces, a workflow engine with dead-letter and replay, and a meeting-autonomy engine that requires human approval — most comparable products implement fragments of this chain. The weakest is that none of it is proven at production load or with paying tenants, so a strategic buyer is paying for accelerated development time, not for proven unit economics.

7. Current Market Benchmarks (2026)

Benchmarks below were retrieved in September 2026 from public M&A; sources; each is a market data point, not an AcquisitionOS-specific fact. Where sources disagree, the disagreement is shown.

Benchmark	2026 figure	Source (retrieved Sep 2026)
Pre-revenue software projects, typical sale	\$500 – \$25,000	exitbid.io, Jul 2026
Micro-SaaS under \$1M ARR, average multiple	~2.85× annual profit (top quartile 6.13×)	Flippa / startupa.ge, Mar 2026
Small SaaS listings, median	2.0× trailing revenue or 3.4× profit (651 listings)	bigideasdb, Jul 2026
Bootstrapped SaaS, typical	~4.8× (revenue basis)	Flippa 2026 guidance
SaaS at \$1M–\$50M ARR	4.0× – 8.0× ARR	AdAstra Equity, Jun 2026
Private SaaS deals, median	~4.7× EV/Revenue (public median ~6.1×)	Aventis Advisors, Aug 2026
Market context	Global TMT deal value +41% YoY to \$2.4T (Jan–May 2026); below ~\$5M EV, EBITDA/SDE multiples dominate	FE International, Aug–Sep 2026

Table 6 — 2026 market benchmarks [MB]. Application to AcquisitionOS starts in Section 8, never as a blind multiple.

Two structural rules from these sources govern this analysis. First, revenue multiples are for businesses with revenue; below roughly \$500K ARR, buyers switch to profit (SDE) multiples, and with zero profit the price collapses toward asset value. Second, pre-revenue assets trade in a narrow band set by code quality, transferability, and buyer-specific strategic fit — which is exactly the Scenario A/B/D structure used next.

8. Valuation Scenarios

Scenario A — Technology / asset sale (as-is today)

Assumes a buyer acquires source code, architecture, AI logic, and product with little or no meaningful recurring revenue — which is the verified state. The pre-revenue benchmark (\$500–\$25K [MB]) sets the frame; AcquisitionOS justifies the upper part of that band and modestly beyond it on depth (505 routes, 258 services, 105 models), coherence, and unusually complete documentation, while hollow e2e tests, 32 failures, the absent license, dev-only database, and unconfigured payments pull the price back down. **Realistic range: \$5,000 – \$60,000 [CE]**, with the low end a fast sale and the upper end requiring a patient, well-marketed process to the right micro-acquirer or agency.

Scenario B — Product acquisition (working SaaS + users)

Assumes a buyer values a working SaaS with users, workflows, integrations, and product maturity. Today this scenario cannot be fully priced because its prerequisite — real users — is unverified; effectively it collapses toward Scenario A. Illustratively [CE]: if the product had 10–30 paying subscribers at \$29–\$99 (ARR ≈ \$6K–\$20K), micro-SaaS profit/revenue multiples of 2–4× would imply \$15K–\$80K before any strategic adjustment. The honest as-is statement: **\$15,000 – \$75,000, conditional on traction materializing that does not currently exist in evidence; as-is, treat as Scenario A.**

Scenario C — Revenue-generating SaaS

With verified MRR of \$0, the revenue-multiple formula yields \$0 — it is included to show the mechanics a buyer would use the moment real revenue exists. Justified multiple for a sub-\$1M-ARR, no-growth-history SaaS: 2–5× ARR [CE, from Table 6 bands].

Hypothetical MRR	ARR	Multiple (2–5×)	Implied EV [CE]	INR at ₹96/USD
\$0 (verified today)	\$0	—	\$0	₹0
\$1,000	\$12,000	2–4×	\$24K – \$48K	₹23L – ₹46L
\$5,000	\$60,000	2.5–4×	\$150K – \$240K	₹1.44Cr – ₹2.3Cr
\$20,000	\$240,000	3–5×	\$720K – \$1.2M	₹6.9Cr – ₹11.5Cr

Table 7 — Scenario C mechanics: what revenue would have to exist for revenue-multiple pricing to matter.

Scenario D — Strategic acquisition

A buyer that can immediately plug the product into an existing customer base and distribution channel can extract value beyond financials: months of roadmap acceleration, a metered-AI reference architecture, and an outreach/automation surface to resell. No invented premium is claimed here; the conditions required for a strategic premium are specific and checkable: (1) the buyer’s roadmap includes lead-gen/SDR capability; (2) integration cost into their stack is low (clean single-process code and docs help); (3) IP is transferable — license file added, founder-ownership clean, no encumbrances; (4) transition support from the founder is agreed (the codebase carries meaningful undocumented context despite good docs); (5) deal structure absorbs risk (earnout or asset warranty). **Conditional range: \$75,000 – \$250,000 [CE]**, with anything above \$100K realistically involving earnout components tied to integration or revenue milestones, not all cash at closing.




Figure 1 — Scenario ranges [CE]. Scenario B is shown at its conditional band; as-is it collapses to Scenario A.

9. Conservative Range

\$5,000 – \$15,000 (₹48,000 – ₹1,44,000 at ₹96/USD) [CE]. This is what a fast, low-risk sale looks like: a micro-acquirer or indie buyer purchasing the codebase as-is, as an asset, with minimal warranties and no transition support beyond a short handover. It is anchored by the pre-revenue benchmark (\$500–\$25K [MB]), positioned in its upper half because the codebase is coherent, documented, and functionally rich rather than a prototype. The assumptions: sale completes within weeks; buyer accepts the product with 32 failing tests, hollow e2e suites, no license file, SQLite-only verified runtime, and unconfigured payment providers; no earnout; founder does minimal knowledge transfer. In a rushed sale to the wrong audience — or with the repository’s fictional metrics document and contradictory legacy docs left in place to poison due diligence — the low end of this band, or no sale at all, is the realistic outcome.

10. Base Range

\$25,000 – \$60,000 (₹24,00,000 – ₹57,60,000) [CE]. This is the most defensible outcome for a reasonably marketed acquisition: a 4–8 week process listing the asset where technical founders, agencies, and micro-SaaS operators look (Acquire.com-style marketplaces, FE International-style brokers for this size, targeted outreach to sales-automation teams). It assumes the pre-sale cleanup in Section 17 is done: license added, failing tests triaged, hollow tests removed or implemented, fictional documents quarantined, a working demo with configured payment sandbox, and a short architecture/DD pack built from the existing docs suite. The upper half of the band requires a buyer who values the documentation and the autonomy architecture; the lower half reflects the median buyer who discounts anything without revenue. Supporting logic: replacement-cost perspective (\$150K–\$400K for the functional core [CE]) tells the buyer they are buying at 15–35 cents on the rebuild dollar — a fair framing that does not pretend rebuild cost equals value.

11. Strategic Range

\$75,000 – \$250,000 (₹72,00,000 – ₹2,40,00,000) [CE], conditional. A strategically aligned buyer — one with an existing customer base, distribution, and a roadmap gap this product fills — could justify paying beyond financial performance for capabilities: the autonomous SDR loop, credits-metered AI architecture, Developer Access API, and a meetings-autonomy engine with human approval. Every dollar of this range depends on the five conditions in Section 8 Scenario D (roadmap fit, low integration cost, clean IP, founder transition support, risk-absorbing structure). Above ~\$100K, expect earnouts or milestone payments rather than all cash: the buyer is paying for integration success they cannot verify pre-close. This range is explicitly not a promise — strategic buyers at this size are rare, processes take months, and the most common failure mode is no strategic buyer appearing at all, in which case the Base Range governs.

12. Key Assumptions and Evidence Classification

Every conclusion in this report carries one of five evidence labels. [V] Verified — read directly from the codebase, database, or configuration during this inspection. [MB] Market benchmark — published 2026 transaction data from named sources (Section 7). [CE] Calculated estimate — derived here from verified inputs and stated logic. [A] Assumption — a condition taken as given for a range. [U] Unknown — no data source exists; the item cannot be established today.

Assumption	Class	If wrong, impact
Only the dev SQLite database is available; no production DB or billing history exists anywhere	V / A	A production DB with real revenue would move pricing to Scenario C mechanics immediately
Buyer acquires code + docs + architecture, not customers/revenue/team	A	Financial buyers would value at near-zero without revenue regardless of code
USD/INR $\approx$ 96 (Sep 21, 2026 observed $\sim$ 95.86)	MB	$\pm$ 5% FX moves INR figures proportionally
Founder provides 2–6 weeks transition support in Base/Strategic cases	A	Without it, subtract roughly 20–30% from buyer willingness [CE]
Pre-sale cleanup (Section 17) is completed before marketing	A	Cleanup skipped $\rightarrow$ Conservative range governs; DD failure risk rises sharply
$\sim$ 348K LOC includes heavy AI-generated volume; replacement cost framed 12–50 engineer-months	CE	A buyer using naive $\text{LOC} \times \text{rate}$ math would overpay; this report refuses that method
No license/ownership encumbrance beyond the missing LICENSE file	A / U	Any co-ownership or third-party code dispute would impair or block a sale

Table 8 — Load-bearing assumptions. Labels follow the Section 12 evidence scheme.

13. Value-Increasing Factors

Only verified strengths are listed; each cites its evidence.

- **End-to-end automation depth** [V]: the full discover $\rightarrow$ score $\rightarrow$ research $\rightarrow$ outreach $\rightarrow$ reply $\rightarrow$ meeting chain exists as working code, not mockups — the rarest object in this buyer class.
- **Four-factor AI scoring with explanations** [V]: reply/conversion/urgency/revenue-potential scores with generated reasoning (scoring-engine.ts, explain-scores routes) — defensible product IP.
- **Credits + entitlements architecture** [V]: per-action metering, enforcement before AI calls, ledger, addons, coupons, two-provider billing scaffolding — monetization plumbing most competitors lack at MVP stage.
- **Developer Access API platform** [V]: hashed keys, rotation, revocation, usage analytics — a genuine platform seam.

- **Security maturity for its size** [V]: MFA, session revocation, device fingerprinting, rate limiting, CSRF/CORS/headers, webhook signature verification, two internal pentest reports with honest findings.
- **Documentation suite** [V]: current (2026-09) architecture docs, ADRs, full API reference, runbooks, KT/handover material — materially lowers technical DD friction and was re-verified and classified during this analysis.
- **Honest self-knowledge** [V]: the feature-status matrix (88 working / 27 partial / 3 planned), scalability plan, and GO/NO-GO review record real limits — buyers price honesty better than surprises.
- **Real unit/integration test core** [V]: 678 passing assertion-backed tests across auth, credits, billing, AI, security subsystems.

14. Value-Reducing Factors

Every factor below is verified in the current codebase; none is speculative.

- **No revenue, no customers** [V]: the single largest value destroyer; removes revenue-multiple pricing entirely.
- **Dev-only runtime** [V]: SQLite sandbox, middleware disabled, no production deployment, no monitoring backend (Sentry/OTLP unset) — buyer inherits a deploy project, not a running business.
- **Payment providers unconfigured** [V]: billing never processed real money; a buyer must validate Stripe/Razorpay flows from scratch.
- **Test integrity gap** [V]: 32 failing tests and ~114 zero-assertion skeleton tests inflate green counts; signals weak release discipline.
- **No LICENSE file** [V]: README badge claims MIT but no license exists — immediate IP transfer blocker.
- **Squashed git history** [V]: 12 commits, initial commit Sep 2026; no development archaeology, no authorship trail — hurts IP provenance confidence.
- **Repo hygiene** [V]: dev.log, db/custom.db, screenshots committed; fictional metrics document and legacy contradictory docs present (now bannerred, but a DD reviewer will still find them).
- **Schema drift** [V]: production schema covers 6 of 105 models; code references fields missing from Prisma models (schema-gap-report.json).
- **Isolation assurance incomplete** [V]: a prior cross-account leakage incident was partially remediated (lead-resolution hardening); without a completed tenancy audit, buyers assume worst-case multi-tenancy risk.
- **Founder/agent concentration** [V]: architecture and operational context live largely with one founder plus AI-agent sessions; no second operator has run the system.
- **Third-party dependencies** [V]: Google Search API, Gmail OAuth, Z-AI/OpenAI, and channel APIs — cost exposure and ToS dependency transfer to the buyer.

15. Buyer Due-Diligence Risks (incl. Founder Dependency & DD Readiness)

A buyer running due diligence today would pass several checks and fail others. The readiness assessment below mirrors a standard technical DD checklist.

DD area	Status today	Buyer concern it triggers
Source organization	Good — coherent src/ layout, thin routes, fat services	None material
Git history	Weak — 12 squashed commits, agent + single human authors	Provenance: who wrote this, and do they have the right to sell it?
Documentation	Strong current suite; legacy contradiction bannered; 1 fictional doc flagged	Which document do I trust? (answered by docs/readme.md + banners)
Architecture docs	Verified modular-monolith classification with diagram + ADRs	None material
Deployment docs	Multi-cloud options documented; verified path = single Railway/Vercel-style web service	Which target is real? (Railway recommendation is authoritative)
Environment/secrets	env-validation + safeguard code; .env present in workspace (must be excluded from transfer)	Secret hygiene; rotated keys required at close
Migrations	3 real migrations + db-push habit; PostgreSQL migration unexecuted	Data-layer maturity; SQLite-to-PG execution risk
Tests	678 real passing; 32 failing; ~114 hollow; no CI config found	Release engineering maturity
Monitoring/logging	In-process logger/OTel hooks; health endpoints; no external backend	Who gets paged when it breaks?
Security	Controls present; pentest reports honest; isolation audit partial	Multi-tenant data leakage is a blocking defect if unfixed
Payments	Coded, idempotent, unconfigured; 2 test invoices	Revenue pipeline unproven
Licenses/dependencies	No LICENSE file; npm tree not audited for copyleft in-report	IP transfer risk — must fix before close
Backup/recovery	scripts/backup + infra/ strategy docs exist	Verify restore actually runs
Account ownership	Sandbox-era accounts (Google, Gmail, Z-AI); no production cloud	Clean handover checklist required (Section 16)

Table 9 — Buyer DD readiness by area.

15.1 Founder dependency

The system is meaningfully founder-dependent despite good documentation. Deployment knowledge (sandbox keepalive scripts, env-restore mechanism, proxy setup), architectural intent behind AI-generated volume, provider account ownership (Google OAuth client, Gmail app password, Z-AI SDK access), and every product decision history sit with one person [V — no second-operator evidence exists in the repo or worklog]. A buyer will price this as transition risk: expect requests for 4–12 weeks of paid transition support, pair-on-call time, or an earnout tied to successful handover. Mitigations that measurably help: record a deployment walkthrough, transfer provider accounts to a dedicated entity, and complete the ops runbook gaps flagged in Section 17.

16. Acquisition Package Checklist

What must be handed over at close, compiled from the verified inventory. Secrets are referenced by location only — never transferred inside the repository.

Category	Items	Verified location / note
Product	Source repo (clean history export), docs/ suite, API reference, KT document	Whole repo minus secrets/runtime artifacts
Infrastructure	Deployment target config (railway.json or chosen platform), Caddyfile/nginx references, keepalive scripts (retire or document)	deploy/railway/; root scripts
Cloud accounts	Hosting account, object storage if used, DNS/domain registrar	Domain ownership: verify — not evidenced in workspace
Database	Schema (105 models), migration history, prod schema variant + its 6-model limitation disclosed, dev DB (test data only)	prisma/
Payments	Stripe/Razorpay accounts + webhook secrets, invoice templates	Accounts must be created/transferred — none exist configured
AI/API accounts	Z-AI or OpenAI keys, Google Search API + CX, SerpAPI if used	.env references — rotate at close
Integrations	Google OAuth client (redirect URIs!), Gmail credentials, Telegram bot, WhatsApp/Twilio if activated	docs/04 + OAUTH_SETUP_GUIDE
Documentation	Architecture + diagram, ADRs, runbooks, security reports, feature-status matrix, this valuation pack	docs/
Legal/IP	LICENSE file (to be created), IP assignment from any contributors, dependency license audit, fictional-docs quarantine record	Gap — required before close
Customer data	None exists [V] — hand over dev DB only as sample data with disclosure	db/custom.db
Analytics	Nothing to hand over (no production analytics) [V]	—
Security	Secrets inventory + rotation record, pentest reports, known-issues list (isolation audit status)	docs/security/, infra/secret-rotation.md
Operations	Cron schedule + CRON_SECRET handling, backup/restore runbooks, monitoring stack (optional, currently not running)	docs/operations/, monitoring/

Table 10 — Handover checklist. All secrets rotate at close; nothing sensitive ships inside the repo.

17. What Must Be Fixed Before Selling (incl. Exit Scorecard)

Ranked by DD-blocking severity. Items 1–5 are effectively mandatory; 6–10 raise the achievable band toward the Base/Strategic ranges.

- **1. Add a LICENSE and clean IP provenance** — decide MIT vs proprietary, record authorship (including AI-assisted authorship disclosure), remove the committed dev DB/dev.log/screenshots. [V gap]
- **2. Complete the account-isolation audit** — finish the 40-resource × all-method tenancy review the product history demands, with automated regression tests; publish the result in the DD pack. [V gap]
- **3. Triage the test suite** — fix or remove the 32 failures; delete or implement the ~114 hollow e2e/load skeletons so green means green. [V gap]
- **4. Quarantine fiction and legacy** — keep the fictional interview document and the abandoned-stack docs out of the DD data room (banners exist; physically separate them for the data room). [V]
- **5. Produce a working demo** — one environment with configured Stripe test mode + Google OAuth + a seeded demo tenant; buyers pay more for what they can click. [A]
- **6. Execute the PostgreSQL migration once** on the full 105-model schema and record the result; retire or reconcile the 6-model production variant. [V gap]
- **7. Stand up minimal external monitoring** — Sentry (free tier) + one health-check uptime monitor; converts "in-process observability" from a claim to evidence. [V gap]
- **8. Record a founder walkthrough** — 60–90 minutes: deploy, env restore, provider accounts, cron, backup/restore. Directly attacks the founder-dependency discount. [A]
- **9. Assemble the DD pack** — architecture summary, this report's Sections 2–5, feature-status matrix, security posture, known-issues list, handover checklist. [A]
- **10. Prepare the improvement story** — the Section 20 table shows buyers what revenue would do to value; it also shows the seller what to build before selling if a fast exit is not forced. [CE]

17.1 Exit readiness scorecard

Category	Current evidence	Missing information	Buyer concern	Required action
Product completeness	88 working / 27 partial features [V]	None material	Partial areas (payments live, CRM sync)	Items 5, 6
Revenue	\$0 verified [V]	Everything financial	Valuation floor	Revenue first, then sell
Growth	No data [U]	Analytics	No momentum signal	Instrument analytics if operating
Retention	No data [U]	Cohorts	Churn unknown	Requires customers first
Technology	Strong depth [V]	Prod DB parity	Migration risk	Item 6
Security	Controls + honest pentests [V]	Completed isolation audit	Tenant leakage	Item 2
Scalability	Limits documented [V]	Load evidence	Single-process ceiling	Queue + PG if scaling story needed
Documentation	Strong [V]	Deploy walkthrough video	Bus factor	Item 8
Operations	Runbooks exist [V]	Live ops evidence	Operational maturity	Item 7
Founder dependency	High [V]	Second operator	Transition risk	Item 8
Customer base	None [V]	—	Priced as asset sale	Revenue first
Financial quality	No financials [U]	Accounting records	Cannot underwrite	Create records if operating
Legal/IP	LICENSE missing [V]	Ownership chain	Transfer blocker	Item 1

Table 11 — Exit readiness scorecard (factual, evidence-cited; no invented composite score).

18. What Additional Data Is Required

To move any conclusion in this report from [U] to established fact — and therefore to re-price with confidence — the following data would need to exist. Each item names the decision it would change.

- **Production database export / billing history** — changes Scenario C from \$0 to a computable EV; the single highest-impact missing artifact.
- **Payment provider statements** (Stripe/Razorpay) — verifies revenue vs subscriptions rows; required for any revenue claim.
- **Analytics export** (signups, WAU/MAU, feature usage) — establishes active users and engagement for Scenario B.

- **Completed isolation audit report** with test results — removes the tenancy-risk discount and the "blocking defect" flag.
- **IP documentation** — contributor assignments, AI-authorship disclosure, dependency license scan — unblocks close.
- **Provider cost history** (AI spend, search API spend) — lets a buyer model gross margin; currently only the metering mechanism exists.
- **Domain and brand assets** — ownership evidence, traffic stats if any; brand is part of the package only if ownership is proven.
- **Churn/retention cohorts** — would upgrade multiple range from generic 2–5× to a defensible band.

19. Final Current Valuation Range

If the entire AcquisitionOS product were put up for sale today — as-is, with the verified evidence in this report — the defensible answer is:

\$5K – \$15K Conservative / fast sale [CE]	\$25K – \$60K Base — most defensible [CE]	\$75K – \$250K Strategic, conditional [CE]
--	---	--

Scenario	USD	INR (at ₹96/USD [MB])	What it assumes
Low / Conservative	\$5,000 – \$15,000	₹48,000 – ₹1,44,000	Fast asset sale, no cleanup, minimal handover, risk-heavy buyer
Base / Most defensible	\$25,000 – \$60,000	₹24,00,000 – ₹57,60,000	4–8 week marketed process, Section 17 items 1–5 done, 2–6 week founder transition
High / Strategic	\$75,000 – \$250,000	₹72,00,000 – ₹2,40,00,000	Aligned buyer with distribution, full cleanup, IP clean, earnout structure likely above \$100K

Table 12 — Final answer. Enterprise value for an asset sale ≈ equity value here (no debt, no cash line to reconcile).

Terms clarification a founder needs at this size: asking price is the number you list; enterprise value and equity value converge for a debt-free asset sale; cash at closing is typically below headline in strategic deals (earnouts, escrows of 10–20%, transition holdbacks); seller financing (a note for part of the price) is common in micro-SaaS and can add 10–30% to achievable price for creditworthy structures [CE/MB]. None of these numbers is a guaranteed sale price; they are the ranges current evidence supports.

20. Potential Valuation After Specific Improvements

Improvements are ranked by valuation mechanics, not effort. Revenue-related rows dominate everything else; no other lever moves value by an order of magnitude. Where reliable evidence does not exist (e.g., precisely how much a specific fix adds), the impact is given as a direction and band with its reasoning — invented point estimates are deliberately omitted.

Metric / improvement	Current state	Target	Potential valuation impact [CE]
MRR (real paying customers)	\$0 verified	\$1K–\$5K MRR	Moves pricing to Scenario C: \$24K–\$240K EV; the only order-of-magnitude lever
ARR sustained + growth	None	\$50K–\$100K ARR, >10% MoM	Enters 3–5× ARR band: \$150K–\$500K
Retention evidence	No cohorts	NRR ≥ 90% documented	Upgrades multiple band ~+0.5–1.0× within same ARR
Customer concentration	N/A	No customer >25% of revenue	Prevents –30% concentration discount
Isolation audit complete	Partial	Full audit + regression tests published	Removes blocking-defect flag; enables strategic range at all
LICENSE + IP chain	Missing	LICENSE + assignments recorded	Unblocks close; without it every range collapses
Test suite integrity	32 fail / ~114 hollow	All green, assertions real, CI configured	Supports upper-Base pricing; weak signal otherwise
Production deployment (PG + monitoring)	Sandbox/SQLite	One live env on Railway/PG + Sentry	Converts "code asset" to "working SaaS" (Scenario B premium)
Founder dependency	High	Walkthrough video + 2nd operator or SOW	Reduces transition discount ~20–30% [CE]
Documentation consistency	Achieved this week [V]	Maintained	Preserves DD speed advantage; low cost, real effect

Table 13 — Improvement → valuation map. Direction/bands only; no invented point estimates.

Closing statement of objectivity: on today's verified evidence — zero revenue, test accounts only, unconfigured payments, incomplete isolation assurance — AcquisitionOS is worth materially less than its build effort and materially more than a weekend prototype. The technology justifies patience; the business metrics justify humility. If a fast sale is desired, sell the asset well after the Section 17 cleanup and expect the Base Range. If revenue is achievable before selling, every month of real MRR does more for the exit number than any other action in this report.